

Time Series Database Interface (TSdbi) Illustrations

Paul D. Gilbert
September 25, 2012

Contents

1	Introduction	1
2	Simple Examples of Time Series Data from the Internet	2
2.1	<i>TShistQuote</i> and <i>TSgetSymbol</i>	2
2.2	<i>TSgetSymbol</i> with FRED	6
2.3	TSxls	11
2.4	TSzip	17
3	SQL Time Series Databases	19
3.1	Writing to SQL Databases	20
4	More Examples	24
4.1	Examples Using SQL Databases	24
4.2	“Inhouse” Database Examples	28
5	Comparing Time Series Databases	30
6	Vintages of Realtime Data	33
7	Appendix A: Connection Specific Details	34
7.1	TSMYSQL Connection Details	35
7.2	TSPostgreSQL Connection Details	36
7.3	TSSQLite Connection Details	37
7.4	TSodbc Connection Details	37
7.4.1	Example ODBC configuration file	38
7.5	TSOracle Connection Details	39
7.6	TSfame Connection Details	40
7.7	TSpadi Connection Details	40
7.8	TS PADI Server	41
8	Appendix B: Underlying Database Structure and Loading Data	41
9	Appendix C: Examples Using DBI and direct SQL Queries	45

1 Introduction

This vignette illustrates the various *R* (R Core Team, 2012) *TSdbi* packages using time series data from several sources. The main purpose of *TSdbi* is to provide a common *API*, so simplicity of changing the data source is a primary feature. The vignette also illustrates some simple time series manipulation and plotting using packages *tframe* and *tfplot*.

To generate this vignette requires most of the *TS** packages, but users will only need one, or a few of these packages. For example, a time series database can be built with several SQL database backends, but typically only one would be used. On the other hand, several packages pull data from various Internet sources, so several of these might be used to accommodate the various specifics of the sources.

If package *TSdata* is installed on your system, it should be possible to view the pdf version of this guide with `vignette("Guide", package="TSdata")`. Otherwise, consider getting the pdf file directly from CRAN at <http://cran.at.r-project.org/web/packages/TSdata/index.html>. Many parts of the vignette can be run by loading the appropriate packages, but some examples use data that has been loaded into a larger database and it will not be possible to reproduce those examples without the underlying database.

Section 2 of this vignette illustrates the mechanism for connecting to a data source. The connection contains all the source specific information so, once the connection is established, the syntax for retrieving data is similar for different sources. The section also illustrates packages that pull data from Internet sources. This currently includes *TSgetSymbol*, *TShistQuote*, *TSxls*, and *TSzip*, but others are in progress. These packages only get data, they do not support writing data to the database. Finally, the section also illustrates the flexibility to return different types of time series objects.

Section 3 illustrates the SQL packages *TSPostgreSQL*, *TSMYSQL*, *TSSQLite*, and *TSodbc*. These use a very standard SQL table structure and syntax, so it should be possible to use other SQL backends. The package *TSOracle* is available on R-forge at <http://tsdbi.r-forge.r-project.org/> but I currently do not have a server in place to test it properly. (If anyone would be interested in doing this, please contact me, pgilbert.ttv9z@ncf.ca.) These packages get data and also support writing data to the database. See Appendix “B” (section 8) and the vignette in package *TSdbi* for more explanation of the underlying database tables, and for bulk loading of data into a database. The section illustrates writing artificial data to the database.

Section 4 provides additional examples of *TSdbi* functionality and of using and graphing series. Subsection 4.1 illustrates fetching data from the web and loading it into a local database. Subsection 4.2 provides examples of how one might use a more extensively populated database that has been previously constructed.

Section 5 illustrates package *TScompare* for comparing time series databases, and Section 6 illustrates the use of realtime vintages of data.

Appendix “A” provides connection details specific to the different database

sources, Appendix “B” provides more details about the structure of SQL databases and, finally, Appendix “C” provides some examples of SQL queries that may be useful for database maintenance.

Many of the *TS** packages are wrappers of other packages. The purpose is to provide a common API for interfacing with time series databases, and an easy mechanism to specify the type of time series object that should be returned, for example, a *ts* object or a *zoo* (Zeileis and Grothendieck, 2005) object. One consequence of providing a common interface is that special strengths of some of the underlying packages cannot always be used. If you really need some of these features then you may need to go directly to the underlying package. However, if you limit your reliance on these underlying features then you will be able to move from one data source to another much more easily.

Loading a *TS** package will also load required packages *TSdbi*, *DBI* (R Special Interest Group on Databases, 2009), *methods*, *tframePlus*, *zoo*, and then any underlying package that the specific *TS** package uses.

2 Simple Examples of Time Series Data from the Internet


This section uses packages to pull data from the Internet. The general syntax of a connection is illustrated, and some simple calculations and graphs are done to demonstrate how the data might be used. However, the purpose of the *TS** packages is to provide a common interface, not to do all time series calculations and graphics. Once you have the data, you should be able to use whatever other *R* packages you like for your calculations and graphs.

The generic aspect of the interface API is accomplished by putting the information specific to the underlying source into the “connection”. Once the connection is established, other aspects of using data are the same, so one connection can be easily interchanged with another, and so your programs do not need to be changed when the data source is changed. (But, of course, some changes will be needed if the names of the variables change.)

2.1 *TShistQuote* and *TSgetSymbol*


Packages *TShistQuote* and *TSgetSymbol* provide mechanisms to retrieve historical quote data from various sources. *TShistQuote* is a wrapper to *get.hist.quote* in package *tseries* (Trapletti and Hornik, 2012). A connection to Yahoo Finance is established, and data retrieved and plotted by

```
> library("TShistQuote")
> yahoo <- TSconnect("histQuote", dbname="yahoo")
> x <- TSget("^gspc", quote = "Close", con=yahoo)
> library("tfplot")
> tfplot(x)
```



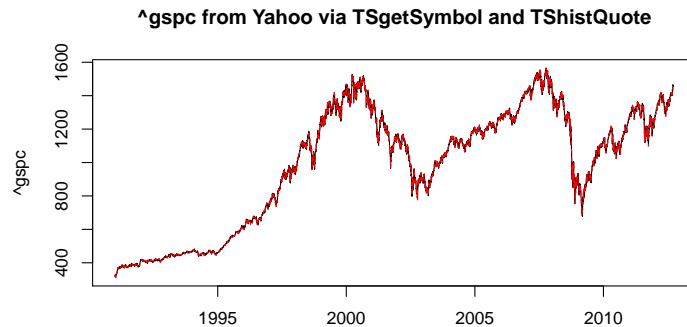
Package *TSgetSymbol* is a wrapper to *getSymbols* in package *quantmod* (Ryan, 2011). A similar connection is established by

```
> library("TSgetSymbol")
> yahoo <- TSconnect("histQuote", dbname="yahoo")
> z <- TSget("^gspc", quote = "Close", con=yahoo)
> tfplot(z)
```



The data from the two connections should be the same:

```
> tfplot(x, z, Title="^gspc from Yahoo via TSgetSymbol and TShistQuote")
```



Here is a second example comparing the same data through two different connection methods:


```
> ya1 <- TSconnect("getSymbol", dbname="yahoo")
> ya2 <- TSconnect("histQuote", dbname="yahoo")
> ibmC1 <- TSget("ibm", ya1, quote = "Close", start="2011-01-03")
> ibmC2 <- TSget("ibm", ya2, quote = "Close", start="2011-01-03")
```

The encoding of the date index vectors are of different classes (Date vs POSIXct), which cause problems comparing these two objects. This can be fixed by:

```
> tframe(ibmC1) <- as.Date(tframe(ibmC1))
```

The two series can then be plotted:

```
> tfplot(ibmC2, ibmC1,
  ylab="IBM Close",
  Title="IBM via getSymbol and histQuote",
  lastObs=TRUE,
  legend=c("via histQuote (black)", "via getSymbol (red)"),
  source="Source: Yahoo")
```



Or the difference can be used to check equality:

```
> max(abs(ibmC2 - ibmC1))
```

```
[1] 0
```

Note that the above does not catch a difference in length, which occurs if new data has been released on one connection and not the other. (At some point Yahoo, was releasing partial data early, and these connections are correcting differently for this. So, at some times of day, the last available data point is not the same on these two connections.)

A certain amount of meta data can be returned with the time series object and can be extracted with these utilities:

```
> TSrefperiod(x)
```

```
[1] "Close"
```

```
> TSdescription(x)
```

```
[1] "^gspc Close from yahoo"
```

```
> TSdoc(x)
```

```
[1] "^gspc Close from yahoo retrieved 2012-09-25 21:09:14"
```

```
> TSlabel(x)
```


```
[1] "^gspc Close"
```

```
> TSsource(x)
```

```
[1] "yahoo"
```

Most of the plots in this vignette will be done with the utilities in the *tfplot* package, but the usual *plot* function produces slightly different results that may be preferable in some situations:

```
> plot(z)
```




Also, for some time series objects, the plot method has been much improved from the default, so if you are using these objects you may find that *plot* provides attractive features.

2.2 *TSgetSymbol* with FRED

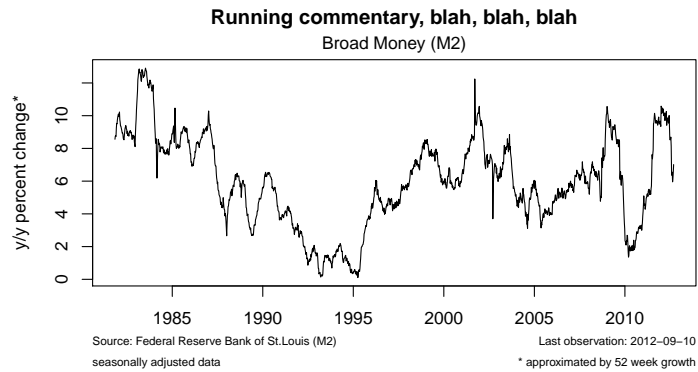
Package *TSgetSymbol* can also be used to get data from the Federal Reserve Bank of St.Louis, as will be illustrated here. (Look at <http://research.stlouisfed.org/fred/>) to find series identifiers.

```
> library("TSgetSymbol")
> fred <- TSconnect("getSymbol", dbname="FRED")
> tfplot(TSget("M2", fred))
```



A connection can be specified to be used as the default, so it does not need to be specified each time:

```
> options(TSconnection=fred)
> tfOnePlot(percentChange(TSget("M2"), lag=52),
  Title = "Running commentary, blah, blah, blah",
  subtitle="Broad Money (M2)",
  ylab= "y/y percent change*",
  source="Source: Federal Reserve Bank of St.Louis (M2)",
  footnoteLeft = "seasonally adjusted data",
  footnoteRight = "* approximated by 52 week growth",
  lastObs = TRUE )
```


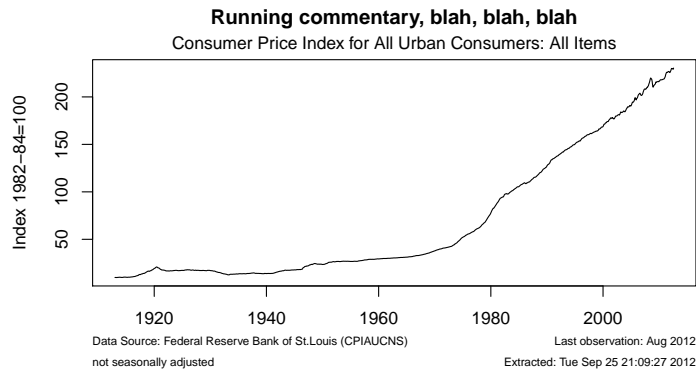


Several different connections will be used in this vignette, and so a default will not be used. To unset the default

```
> options(TSconnection=NULL)
```

It is also possible to return multiple series, but they should all be of the same frequency. (The FRED series called M2 is a weekly series).

```
> x <- TSget(c("CPIAUCNS", "M2SL"), fred)
> tfplot(x,
  Title = "Running commentary, blah, blah, blah",
  subtitle=c("Consumer Price Index for All Urban Consumers: All Items", "Broad Money"),
  ylab= c("Index 1982-84=100", "Billions of dollars"),
  source= c("Data Source: Federal Reserve Bank of St.Louis (CPIAUCNS)",
    "Data Source: Federal Reserve Bank of St.Louis (M2SL)"),
  footnoteLeft = c("not seasonally adjusted", "seasonally adjusted"),
  footnoteRight = paste("Extracted:", date()),
  lastObs = TRUE )
> TSdescription(x)
```

```
> TSdates(c("CPIAUCNS","M2SL"), fred) # note default start
```

```
      [,1]
[1,] "CPIAUCNS from 1913 1 to 2012 8      12"
[2,] "M2SL from 1959 1 to 2012 8      12"
```

By default, *TSget* returns *ts* time series for annual, quarterly, and monthly data, and *zoo* series otherwise. It is possible to specify the type of object to return:

```
> x <- TSget(c("CPIAUCNS","M2SL"), fred, TSrepresentation="zoo")
> class(x)
```

A session default can also be set for this with

```
> options(TSrepresentation="zoo")
```


in which case all results will be return as *zoo* objects unless otherwise specified. The session default is unset with

```
> options(TSrepresentation=NULL)
```

Beware that it does not make sense to set *ts* as the default, because it is already the default for all series that can be represented as *ts*, and will not work correctly for other series. Other representations are possible. See the *TSget* help for more details.

The following connects to yahoo and loads the ticker symbol for Ford. This is a multivariate time series with open, close, etc.

```
> yahoo <- TSconnect("getSymbol", dbname="yahoo")
> x <- TSget("F", con=yahoo)
> plot(x)
```



In this case, *tfplot* displays graphs in verticle panels. However, six panels do not fit on a page, so only the first three are displayed with:


```
> tfplot(x,series=1:3)
```



It is possible to specify the number of graphs on a screen with *graphs.per.page*, for example, `tfplot(x, graphs.per.page=3)`. Set `par(ask=TRUE)` if you want to stop and prompt for <Return> between pages in the graphics output.

The *quote* argument can be used to specify that only a subset of the market data should be returned:

```
> tfOnePlot(TSget("F", con=yahoo, quote=c("Open", "Close")),
  Title="Ford from Yahoo; Open (black); Close (red)",
  ylab="Price")
```



2.3 TSxls

TSxls provides methods for the *TSdbi* interface, allowing the use of spreadsheets as if they are a database. (This is a poor substitute for a real database, but is sometimes convenient.) *TSxls* use package *read.xls* in *gdata* (Warnes et al., 2011) *TSxls* does not support writing data to the spreadsheet (but see *TSwriteXLS* in *tframePlus* to write time series data into a spreadsheet). The spreadsheet can be a remote file, which is retrieved when the connection is established.

The following retrieves a file from the Reserve Bank of Australia and maps the elements that are used: data, dates, identifiers, and series names.

```
> library("TSxls")
> rba <- TSconnect("xls",
  dbname="http://www.rba.gov.au/statistics/tables/xls/d03hist.xls",
  map=list(ids =list(i=11, j="B:Q"),
    data =list(i=12:627, j="B:Q"),
    dates=list(i=12:627, j="A"),
    names=list(i=4:7, j="B:Q"),
    description = NULL,
    tsrepresentation = function(data,dates){
      ts(data,start=c(1959,7), frequency=12)}))
```


This also illustrates how *tsrepresentation* can be specified as an arbitrary function to set the returned time series object representation.

Beware that data is read into *R* when the connection is established, so changes in the spreadsheet will not be visible in *R* until a new connection is established.

```

> x <- TSget("DMACN", rba)
> require("tfplot")
> tfplot(x)

```





```

> x <- TSget(c("DMAM1N", "DMAM3N"), rba)
> tfplot(x)
> TSdescription(x)

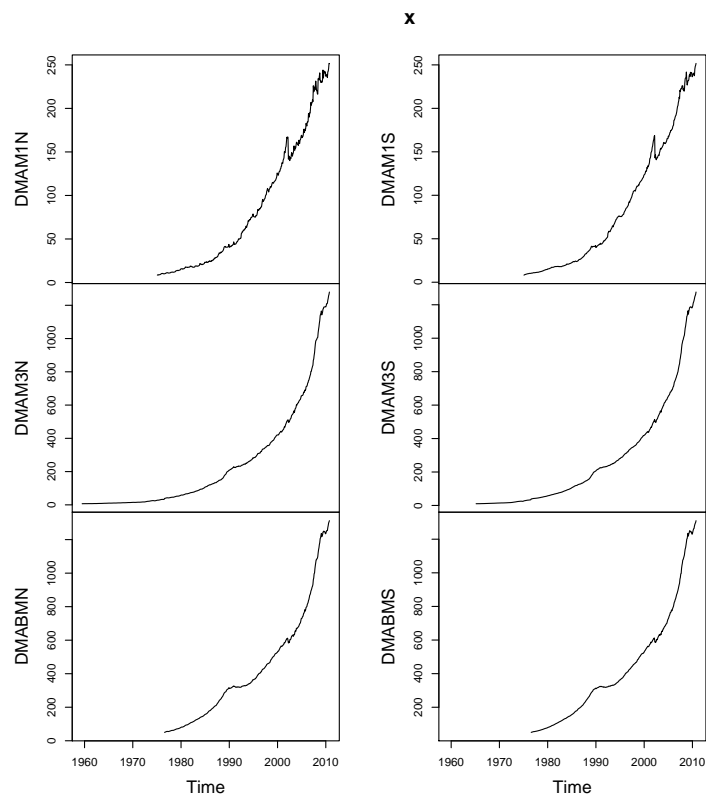
[1] " from http://www.rba.gov.au/statistics/tables/xls/d03hist.xls"
[2] " from http://www.rba.gov.au/statistics/tables/xls/d03hist.xls"

```





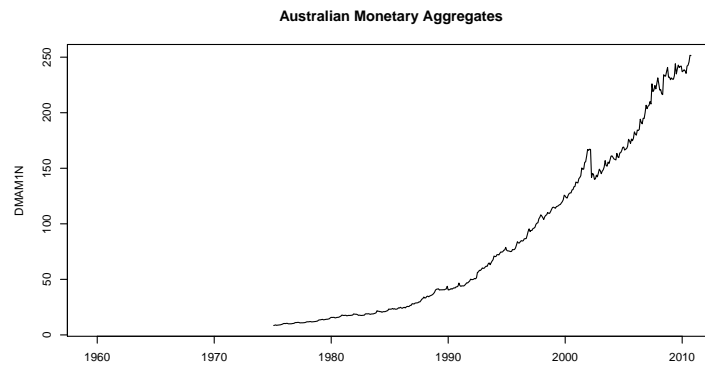
As already illustrated above, the function *plot* displays the series somewhat differently than *tfplot*, and possibly differently depending on the objects time series representation. *tfplot* treats each series in the first argument as panels to be plotted.

```
> x <- TSget(c("DMAM1N", "DMAM3N", "DMABMN", "DMAM1S", "DMAM3S", "DMABMS"), rba)
> plot(x)
```






tfplot treats each series in the first argument as a panel to be plotted. It is possible to specify the number of graphs on each page of the output device with the argument *graphs.per.page*. It is also possible to specify that a subset of the series should be selected:

```
> tfplot(x, Title="Australian Monetary Aggregates", series=1:3)
```




```
> tfplot(x, Title="Australian Monetary Aggregates", series=4:6)
```

tfplot takes additional time series objects as arguments. Series in the first argument are plotted in separate panels. Series in subsequent time series objects will be plotted respectively on the same panels as the first, so the number of series in each object must be the same.

```
> tfplot(TSget(c("DMAM1S", "DMAM3S", "DMABMS"), rba),
         TSget(c("DMAM1N", "DMAM3N", "DMABMN"), rba),
```

```
ylab=c("DMAM1", "DMAM3", "DMABM"),
Title="Australian Monetary Aggregates")
```





2.4 TSzip

TSzip provides methods for the *TSdbi* interface, allowing the use of zipped files that can be read by *read.table* as if each file is a database series (or group of series such as high, low, open, close, for a stock). The *dbname* is a directory or *url*. *TSzip* does not support writing data to the database.


The following retrieves zipped files from http://pittrading.com/free_market_data.htm which provides some end of day data free of charge. (Disclaimer: This site is used as an example. Other than using this free data, I have no association with the company.)

```
> library("TSzip")
> pitr <- TSconnect("zip", dbname="http://pittrading.com/free_eod_data")
> z <- TSget("INDU", pitr)
> tfplot(z, series=c(1,4))
```



The following illustrates returning an *xts* (Ryan and Ulrich, 2011) time series object.

```
> z <- TSget(c("EURUSD", "GBPUSD"), pitr, quote=c("Open","Close"),
             TSrepresentation=xts)
> tfplot(z,
          Title="EURUSD and GBPUSD open and closing values from pitrading",
          start="1995-01-01",
          par=list(omi=c(0.1,0.3,0.1,0.1),mar=c(2.1,3.1,1.0,0.1)))
```



The default behaviour of printed graphs above is improved by adjusting graphics device margins. (They are set by the vector in order: bottom, left, top, right.) It is often useful to adjust these when generating *pdf* files for publication.

3 SQL Time Series Databases

This section gives several simple examples of putting series on and reading them from a database. (If a large number of series are to be loaded into a database, one would typically do this with a batch process using the database program's utilities for loading data.) The examples in this section will use *TSMYSQL* but, other than the initial connection, access will be similar for other SQL *TS** packages. The syntax for connecting with other packages, and other options for connecting with *TSMYSQL*, are provided in Appendix "A".

The packages *TSPostgreSQL*, *TSMYSQL*, *TSSQLite*, and *TSodbc* use underlying packages *RPostgreSQL* (Conway et al., 2012), *RMySQL* (James and DebRoy, 2012), *RSQLite* (James, 2011), and *RODBC* (Ripley and from 1999~to Oct 2002 Michael Lapsley, 2012) to provide access to an SQL database with an underlying table structure which is set up to store time series data.

The next lines of code do some preliminary setup of the database. This uses the underlying database connection (*dbConnect*) rather than *TSconnect*, because *TSconnect* will not recognize the database until it has been setup.

WARNING: running this will overwrite the “test” database on your server.

```
> library("TSMysql")
> con <- dbConnect("MySQL", dbname="test")
> source(system.file("TSql/CreateTables.TSql", package = "TSdbi"))
> dbDisconnect(con)
```

3.1 Writing to SQL Databases

This subsection illustrates writing some simple artificial data to a database, and reading it back. This part of the vignette is generated using *TSMysql*, but other backend SQL servers work in a similar way. See Appendix “A” for details of establishing other SQL database connections.

The first thing to do is to establish a *TSdbi* connection to the database:

```
> library("TSMysql")
> con <- TSconnect("MySQL", dbname="test")
```

TSconnect uses *dbConnect* from the *DBI* package, but checks that the database has expected tables, and checks for additional features. (It cannot be used before the tables are created, as was done above.)

The follow illustrates the use of the *TSdbi* interface, which is common to all extension packages.

This puts a series called *vec* on the database and then reads it back

```
> z <- ts(rnorm(10), start=c(1990,1), frequency=1)
> seriesNames(z) <- "vec"
> if(TSexists("vec", con)) TSdelete("vec", con)
> TSput(z, con)
> z <- TSget("vec", con)
```

Note that the series name(s) and not the *R* variable name (in this case, *vec* not *z*) are used on the database. If the retrieved series is printed it is seen to be a “ts” time series with some extra attributes.



TSput fails if the series already exists on the *con*, so the above example checks and deletes the series if it already exists. *TSreplace* does not fail if the series does not yet exist, so examples below use it instead. Several plots below show original data and the data retrieved after it is written to the database. In the plot below, one is added to the original data so that both lines are visible.

The *R* variable can contain multiple series of the same frequency. They are stored separately on the database.

```
> z <- ts(matrix(rnorm(200),100,2), start=c(1995,1), frequency=12)
> seriesNames(z) <- c("mat2c1", "mat2c2")
> TSreplace(z, con)
```

```
[1] TRUE
```

```
> tfplot(z+1, TSget(c("mat2c1","mat2c2"), con),  
          lty=c("solid", "dashed"), col=c("black", "red"))
```



The following extract information about the series from the database, although not much information has been added for these examples.

```
> TSmeta("mat2c1", con)  
  
serIDs: mat2c1  
from dbname test using TMySQLConnection  
  
> TSmeta("vec", con)  
  
serIDs: vec  
from dbname test using TMySQLConnection  
  
> TSdates("vec", con)  
  
[ ,1]  
[1,] "vec from 1990 1 to 1999 1 A      "
```

```
> TSdescription("vec", con)
```

```
[1] NA
```

```
> TSdoc("vec", con)
```

```
[1] NA
```

```
> TSlabel("vec", con)
```

```
[1] NA
```

Data documentation can be in three forms. A description specified by *TSdescription*, longer documentation specified by *TSdoc*, or a short label, typically useful on a graph, specified by *TSlabel*. These can be added to the time series object, in which case they will be written to the database when *TSput* or *TSreplace* is used to put the series on the database. Alternatively, they can be specified as arguments to *TSput* or *TSreplace*. The description or documentation will be retrieved as part of the series object with *TSget* only if this is specified with the logical arguments *TSdescription* and *TSdoc*. They can also be retrieved directly from the database with the functions *TSdescription* and *TSdoc*.

```
> z <- ts(matrix(rnorm(10),10,1), start=c(1990,1), frequency=1)
> TSreplace(z, serIDs="Series1", con)
```

```
[1] TRUE
```

```
> zz <- TSget("Series1", con)
> TSreplace(z, serIDs="Series1", con,
  TSdescription="short rnorm series",
  TSdoc="Series created as an example in the vignette.")
```

```
[1] TRUE
```

```
> zz <- TSget("Series1", con, TSdescription=TRUE, TSdoc=TRUE)
> start(zz)
```

```
[1] 1990    1
```

```
> end(zz)
```

```
[1] 1999    1
```

```
> TSdescription(zz)
```

```
[1] "short rnorm series"
```

```
> TSdoc(zz)
```

```

[1] "Series created as an example in the vignette."

> TSdescription("Series1", con)

[1] "short rnorm series"

> TSdoc("Series1", con)

[1] "Series created as an example in the vignette."

```

The following examples use dates and times which are not handled by *ts*, so the *zoo* time representation is used. It is necessary to specify the table where the data should be stored in cases where it is difficult to determine the periodicity of the data. See Appendix “B” for details of the specific tables.



```

> require("zoo")
> z <- zoo(matrix(rnorm(200),100,2), as.Date("1990-01-01") + 0:99)
> seriesNames(z) <- c("zooc1", "zooc2")
> TSreplace(z, con, Table="D")

[1] TRUE

> tfplot(z+1, TSget(c("zooc1", "zooc2"), con),
          lty=c("solid", "dashed"), col=c("black", "red"))
>

```

```
> z <- zoo(matrix(rnorm(200),100,2), as.Date("1990-01-01") + 0:99 * 7)
> seriesNames(z) <- c("zooWc1", "zooWc2")
> TSreplace(z, con, Table="W")

[1] TRUE

> dbDisconnect(con)
> #dbUnloadDriver(m)
> detach(package:TMySQL)
> detach(package:RMySQL)
```

4 More Examples

4.1 Examples Using SQL Databases

This section illustrates fetching data from the web and loading it into the database. This would be a very slow way to load a database, but provides examples of different kinds of time series data. The fetching is done with *TShistQuote*.

Fetching data can fail due to lack of an Internet connection or delays, which will cause the generation of this vignette to fail.

This part of the vignette is generated using *TSPostgreSQL*, but other back-end SQL servers work in a similar way. See Appendix “A” for details of establishing other SQL database connections.

First establish a connection to the database where data will be saved:

```
> require("TSPostgreSQL")
> host <- Sys.getenv("POSTGRES_HOST")
> con <- TSconnect("PostgreSQL", dbname="test", host=host)
```

Now connect to the web server and fetch data:

```
> require("TShistQuote")
> yahoo <- TSconnect("histQuote", dbname="yahoo")
> x <- TSget("^gspc", quote = "Close", con=yahoo)
```

Then write the data to the local server, specifying table B for business day data (using *TSreplace* in case the series is already there from running this example previously):

```
> TSreplace(x, serIDs="gspc", Table="B", con=con)
```

```
[1] TRUE
```

and check the saved version:

```
> TSrefperiod(TSget(serIDs="gspc", con=con))
```

```
[1] "Close"
```

```
> TSdescription("gspc", con=con)
```

```
[1] "^gspc Close from yahoo"
```


```
> TSdoc("gspc", con=con)
```

```
[1] "^gspc Close from yahoo retrieved 2012-09-25 21:11:44"
```

```
> TSlabel("gspc", con=con) # this may not yet be supported on the db
```

```
[1] NA
```

```
> tfplot(TSget(serIDs="gspc", con=con))
```



```

> x <- TSget("ibm", quote = c("Close", "Vol"), con=yahoo)
> TSreplace(x, serIDs=c("ibm.Cl", "ibm.Vol"), con=con, Table="B",
  TSdescription.=c("IBM Close","IBM Volume"),
  TSdoc.= paste(c(
    "IBM Close retrieved on ",
    "IBM Volume retrieved on "), Sys.Date()))

[1] TRUE

> z <- TSget(serIDs=c("ibm.Cl", "ibm.Vol"),
  TSdescription=TRUE, TSdoc=TRUE, con=con)
> TSdescription(z)



[1] "IBM Close" "IBM Volume"

> TSdoc(z)

[1] "IBM Close retrieved on 2012-09-25"
[2] "IBM Volume retrieved on 2012-09-25"

> tfplot(z, xlab = TSdoc(z), Title = TSdescription(z))
> tfplot(z, Title="IBM", start="2007-01-01")

```



Oanda has maximum of 500 days, so the start date is specified here so as to not exceed that.

```
> Oanda <- TSconnect("histQuote", dbname="oanda")
> x <- TSget("EUR/USD", start=Sys.Date() - 495, con=Oanda)
> TSreplace(x, serIDs="EUR/USD", Table="D", con=con)
```

```
[1] TRUE
```

Then check the saved version:

```
> z <- TSget(serIDs="EUR/USD", TSlabel=TRUE, TSdescription=TRUE, con=con)
> tfplot(z, Title = TSdescription(z), ylab=TSlabel(z), start="2007-03-01")
```



```
> dbDisconnect(con)
> dbDisconnect(yahoo)
> dbDisconnect(Oanda)
> detach(package:TSPostgreSQL)
> detach(package:RPostgreSQL)
```

4.2 “Inhouse” Database Examples

These examples use a database which has been called “ets”. It has already been populated with some series from various sources, so the examples illustrate how a large “inhouse” database might be used. Typically a large database would be installed using database scripts directly rather than from *R* with *TSpout* or *TSreplace*.

The following uses *TSMYSQL* but, other than the initial connection, everything should remain the same if another database is used. (See Appendix A regarding the details of other connections.)

First, establish a connection to the database, and make it the default.

```
> library("TSMYSQL")
> ets <- TSconnect("MySQL", dbname="ets")
> options(TSconnection=ets)
```

Often one needs to check if data is available or has been updated. This defines a list of series ids, with more explanatory names, and checks availability:

```
> EXCH.IDs <- matrix(c(
  "M.SDR.CCUSMA02.ST",      "SDR/USD exchange rate",
  "M.CAN.CCUSMA02.ST",      "CAN/USD exchange rate",
  "M.MEX.CCUSMA02.ST",      "MEX/USD exchange rate",
  "M.JPN.CCUSMA02.ST",      "JPN/USD exchange rate",
  "M.EMU.CCUSMA02.ST",      "Euro/USD exchange rate",
  "M.OTO.CCUSMA02.ST",      "OECD /USD exchange rate",
  "M.G7M.CCUSMA02.ST",      "G7   /USD exchange rate",
```

```

      "M.E15.CCUSMA02.ST",      "Euro 15. /USD exchange rate"
    ), 2, 8)
> TSdates(EXCH.IDs[1,])

[,1]
[1,] "M.SDR.CCUSMA02.ST from 1960 1 to 2009 2 M NA  "
[2,] "M.CAN.CCUSMA02.ST from 1960 1 to 2009 2 M NA  "
[3,] "M.MEX.CCUSMA02.ST from 1963 1 to 2009 2 M NA  "
[4,] "M.JPN.CCUSMA02.ST from 1960 1 to 2009 2 M NA  "
[5,] "M.EMU.CCUSMA02.ST from 1979 1 to 2009 2 M NA  "
[6,] "M.OTO.CCUSMA02.ST not available"
[7,] "M.G7M.CCUSMA02.ST not available"
[8,] "M.E15.CCUSMA02.ST not available"


```

(The available data in this test data set is fairly old.)
The following examples illustrate using metadata in the graph.

```

> z <- TSget("V122646", TSdescription=TRUE, TSSource=TRUE, TSdoc=TRUE,
  TSlabel=TRUE )
> tfplot(z, Title=strsplit(TSdescription(z), ","))

```



Some metadata fields in my example database are empty. If they were available, this step would not be necessary. The missing meta data can be put into the returned object by:

```

> TSSource(z) <- "Bank of Canada"
> TSlabel(z) <- "Short-Term Business Credit"


```

With this, more of the meta data can be indicated on the graph:

```

> tfplot(ytoypc(z), ylab="Year-over-Year Growth",
  Title=TSlabel(z),
  source=paste(TSseriesIDs(z), "from", TSSource(z)),
  lastObs=TRUE,
  footnote=TSdescription(z))

```



```
> dbDisconnect(ets)
> options(TSconnection=NULL)

> detach(package:TMySQL)
> detach(package:RMySQL)
```

5 Comparing Time Series Databases

The purpose of package *TScompare* is to compare pairs of series on two database. These series might have the same name, but for generality the main function, *TScompare*, is set up to use name pairs. The pairs to compare are indicated by a matrix of strings with two columns. Alternately, it would be possible to compare pairs on the same database but, other than for testing or demonstration purposes, this would not make sense unless the names are different.

The connections are established using other *TSdbi* packages such as *TMySQL*, *pkgTSPostgreSQL*, etc. It will be necessary to establish two database connections, so it will also be necessary to load the database specific packages. In this vignette, examples will use *TShistQuote*, *TMySQL* and *TSSQLite*.

```
> library("TScompare")
> library("TShistQuote")
> library("TMySQL")
> library("TSSQLite")
```

First setup database tables that are used by *TSdbi* using a *dbConnect* connection, after which a *TSconnect* connection can be used (This is re-done to insure the database are empty):

```
> con <- dbConnect("MySQL", dbname="test")
> source(system.file("TSql/CreateTables.TSql", package = "TSdbi"))
> dbDisconnect(con)
> con <- dbConnect("SQLite", dbname="test")
```

```
> source(system.file("TSql/CreateTables.TSql", package = "TSdbi"))
> dbDisconnect(con)
```

Now TS connections to the databases are established.

```
> con1 <- TSconnect("MySQL", dbname="test")
> con2 <- TSconnect("SQLite", dbname="test")
```

Next a connection to yahoo is used to get some series and write them to the local test database. (See the vignette for *TShistQuote* for more examples of reading series from the web.) *TSreplace* is used because *TSput* will fail if the series already exists.

```
> yahoo <- TSconnect("histQuote", dbname="yahoo")
> x <- TSget("~ftse", yahoo)
> TSreplace(x, serIDs="ftse", Table="B", con=con1)

[1] TRUE

> TSreplace(x, serIDs="ftse", Table="B", con=con2)

[1] TRUE

> x <- TSget("~gspc", yahoo)
> TSreplace(x, serIDs="gspc", Table="B", con=con1)

[1] TRUE

> TSreplace(x, serIDs="gspc", Table="B", con=con2)

[1] TRUE

> x <- TSget("ibm", con=yahoo, quote = c("Close", "Vol"))
> TSreplace(x, serIDs=c("ibmClose", "ibmVol"), Table="B", con=con1)

[1] TRUE

> TSreplace(x, serIDs=c("ibmC", "ibmV"), Table="B", con=con2)

[1] TRUE
```

Now to do a comparison:

```
> ids <- AllIds(con1)
> ids

[1] "ftse"      "gspc"      "ibmClose" "ibmVol"
```

If the second database has the same names then ids can be made into a matrix with identical columns.


```

> ids <- cbind(ids, ids)
> eq <- TScompare(ids, con1, con2, na.rm=FALSE)
> summary(eq)

4 of 4 are available on con1.
2 of 4 are available on con2.
2 of 2 remaining have the same window.
2 of 2 remaining have the same window and values.

> eqrm <- TScompare(ids, con1, con2, na.rm=TRUE)
> summary(eqrm)

4 of 4 are available on con1.
2 of 4 are available on con2.
2 of 2 remaining have the same window.
2 of 2 remaining have the same window and values.

```

Since names are not identical the above indicates discrepancies, which are resolved by indicating the corresponding name pairs:

```

> ids <- matrix(c("ftse","gspc","ibmClose", "ibmVol",
                  "ftse","gspc","ibmC", "ibmV"),4,2)
> ids

      [,1]      [,2]
[1,] "ftse"    "ftse"
[2,] "gspc"    "gspc"
[3,] "ibmClose" "ibmC"
[4,] "ibmVol"   "ibmV"

> eq <- TScompare(ids, con1, con2, na.rm=FALSE)
> summary(eq)

4 of 4 are available on con1.
4 of 4 are available on con2.
4 of 4 remaining have the same window.
4 of 4 remaining have the same window and values.

> eqrm <- TScompare(ids, con1, con2, na.rm=TRUE)
> summary(eqrm)

4 of 4 are available on con1.
4 of 4 are available on con2.
4 of 4 remaining have the same window.
4 of 4 remaining have the same window and values.

```

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```

> detach(package:TMySQL)
> detach(package:RMySQL)
> detach(package:TSSQLite)
> detach(package:RSQLite)

```

6 Vintages of Realtime Data

Data vintages, or “realtime data” are snapshots of data that was available at different points in time. The most obvious feature of earlier snapshots is that the series end earlier. However, the reason for retaining vintages is that data is often revised so, for some observations, earlier vintages have different data. Typically most revisions happen for the most recent periods, but this is often the data of most interest for forecasting and policy decisions, so the revision records are valuable for understanding the implications of differences between early and revised releases of the data. A simple mechanism for accessing vintages of data is available in several *TS** packages. This is illustrated here with an SQL database that has been set up with vintage support.

First establish a connection to the database and get a vector of the available vintages:

```

> require("TMySQL")
> require("tfplot")
> ets <- TSconnect("MySQL", dbname="etsv")
> v <- TSvintages(ets)

```

The following checks if a certain variable (Consumer Credit – V122707) is available in the different vintages. (V numbers replaced B numbers circa 2003, so the V numbers do not exist in older vintages. This could be supported by implementing aliases, but that has not been done here.)

```

> ve <- TSexists("V122707", vintage=v, con=ets)

```

A temporary problem with the db loading of some vintages is fixed by

```

> ve[224:length(ve)] <- FALSE


```

The vintages can then be retrieved and plotted by

```

> CC <- TSget(serIDs="V122707", con=ets, vintage=v[ve])#BUG this is throwing a warning
> tfOnePlot(ytoypc(CC), start=c(2000,1),
  ylab="Consumer Credit (V122707) y/y Growth",
  Title=paste("Vintages", v[ve][1], "to", v[ve][189]),
  lastObs=TRUE, source="Source: Bank of Canada")

```



With package *TSfame* vintages are supported if the vintages are stored in files with names like “etsmfacansim_20110513.db”. Then the vintages can be accessed as follows:

```
> dbs <- paste("ets /path/to/etsmfacansim_", c(
  "20110513.db", "20060526.db", "20110520.db"), sep="")
> names(dbs) <- c("2011-05-13", "2006-05-26", "2011-05-20")
> conetsV <- TSconnect("fame", dbname=dbs, "read", current="2011-05-13")
> z <- TSget("V122646", con=conetsV, vintage=c("2011-05-13", "2006-05-26"))
> dbDisconnect(conetsV)
```

(The above example should work, but beware that I am no longer testing it because I no longer have Fame access.)

The package *googleVis* can be used to produce a plot that is very useful for examining vintage more closely, and finding outliers and other data problems. The names are used in the legend of this next plot, so set series names. (Otherwise they get reset to y/y is messy to read.)

```
> require("googleVis")
> tfVisPlot(ytoypc(CC, names=seriesNames(CC)), start=c(2006,1),
  options=list(title="Vintages of Consumer Credit (V122707) y/y Growth"))
```

This will produce a graph in your web browser. It is not reproduced here. (And beware that it may not be very fast.) Pointing your mouse at the legend of this plot will highlight that vintage, and pointing at the graph will give information about the source of a data point.

7 Appendix A: Connection Specific Details

This section provides details of the different connections which are specific to individual packages and backend databases. In order to make the examples complete, for the SQL versions, test databases are first created with the tables

expected by the *TS** packages. Note that this is done with a *dbConnect* connection rather than a *TSconnect* connection, because *TSconnect* expects the tables to exist already.

WARNING: running these example will overwrite tables in the “test” database on the server.

The database setup might typically be done by an administrator, rather than by an end user. Here it is done using a small script *CreateTables.TSsql* distributed with the *TSdbi* package. A more detailed description of the instructions for building the database tables is given in the vignette for the *TSdbi* package. Those instruction show how to build the database using database utilites rather than *R*, which might be the way a system administrator would build the database.

In many cases there are two or more ways to pass information like the *username*, *password*, and *server* or *host*. One mechanism is that this information is specified in a configuration file in the user’s home directory. The database driver then reads this information and it is not part of the user’s *R* session. (Often this is considered the most secure way.) Another way is that environment variables are set, and the database driver uses these. Again, this is not part of the user’s *R* session. Still another way is that the user passes this information in the call to *TSconnect* in their *R* session. In this case the character strings are visible in the *R* session, and possibly recorded in the user’s *R* scripts, thus this is typically not considered to be very secure. A modification, which is only a little bit better, is for the user’s *R* scripts to read the information from environment variables using, for example:

```
> user <- Sys.getenv("MYSQL_USER")
```

7.1 TSMYSQL Connection Details

The MySQL user, password, and hostname should be set in MySQL client configuration file (*.my.cnf*) in the user’s home directory before starting *R*. Alternatively, this information can be set with environment variables *MYSQL_USER*, *MYSQL_PASSWD* and *MYSQL_HOST*. (An environment variable *MYSQL_DATABASE* can also be set, but “test” is specified below.) Below the configuration file is used.

The next small section of code uses *dbConnect* to set up database tables that expected by *TSconnect*.

```
> library("TSMYSQL")
> con <- dbConnect("MySQL", dbname="test")
> source(system.file("TSsql/CreateTables.TSsql", package = "TSdbi"))
> dbDisconnect(con)
```

Now a *TSdbi* connection to the database is established.

```
> con <- TSconnect("MySQL", dbname="test")
```

Another alternative is to pass the user/password information in the arguments to the connection function:

```
> con <- TScconnect("MySQL", dbname="test", username=user, password=passwd, host=host)
```

This is may be cumbersome to change, and is generally considered to be less secure.

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```
> detach(package:TMySQL)
> detach(package:RMySQL)
```

7.2 TSPostgreSQL Connection Details

The PostgreSQL user, and password, can be set in PostgreSQL configuration file (.pgpass in Linux) in the user's home directory before starting R. The Postgress documentation suggests that it should be possible to get the host from the .pgpass file too, but I have not been able to make that work. The PostgreSQL alternative to the configuration file is to use environment variables PGDATABASE, PGHOST, PGPORT, and PGUSER. This package supports another alternative to set this information with environment variables POSTGRES_USER, POSTGRES_PASSWD and POSTGRES_HOST, which are read in the R code. (An environment variable POSTGRES_DATABASE can also be set, but "test" is specified below.) Below, the environment variable POSTGRES_HOST is used to determine the host server, but the .pgpass file is used for the user and password information.

```
> host <- Sys.getenv("POSTGRES_HOST")
```

The next small section of code uses *dbConnect* to set up database tables that expected by *TScconnect*.

```
> library("TSPostgreSQL")
> con <- dbConnect("PostgreSQL", dbname="test", host=host)
> source(system.file("TSql/CreateTables.TSql", package = "TSdbi"))
> dbDisconnect(con)
```

Now a TSdbi connection to the database is established.

```
> con <- TScconnect("PostgreSQL", dbname="test", host=host)
```

Another alternative is to pass the user/password information in the arguments to the connection function:

```
> con <- TScconnect("PostgreSQL", dbname="test", user=user, password=passwd, host=host)
```

This is may be cumbersome to change, and is generally considered to be less secure.

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```
> detach(package:TSPostgreSQL)
> detach(package:RPostgreSQL)
```

7.3 TSSQLite Connection Details

In SQLite there does not seem to be any need to set user or password information, and examples here all use the localhost.

Now setup database tables that are used by TSdbi using a *dbConnect* connection, after which a *TScconnect* connection can be used:

```
> library("TSSQLite")
> con <- dbConnect("SQLite", dbname="test")
> source(system.file("TSql/CreateTables.TSql", package = "TSdbi"))
> dbDisconnect(con)
```

Now a TSdbi connection to the database is established.

```
> con <- TScconnect("SQLite", dbname="test")
```

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```
> detach(package:TSSQLite)
> detach(package:RSQLite)
```

7.4 TSodbc Connection Details

The ODBC user, password, hostname, etc, should be set in ODBC client configuration file in the user's home directory (.odbc.ini in Linux) before starting R. An example of this file is provided below. It will also be necessary to have the appropriate driver installed on the system (Postgresql in the example below). Alternatively, it should be possible to set this information with environment variables ODBC_USER, ODBC_PASSWD and ODBC_DATABASE. However, the variable ODBC_HOST does not seem to work for passing the ODBC connection, so a properly setup ODBC configuration file is needed. Because of this, the environment variable mechanism is not currently supported in *TSodbc* and the user, passwd, and host settings should be done in the configuration file.

Now setup database tables that are used by TSdbi using a *odbcConnect* connection, after which a *TScconnect* connection can be used:

```
> library("TSodbc")
> con <- odbcConnect(dsn="test")
> if(con == -1) stop("error establishing ODBC connection.")
> source(system.file("TSql/CreateTables.TSql", package = "TSdbi"))
> odbcClose(channel=con)
```

Now a TSdbi connection to the database is established.

```
> con <- TScconnect("ODBC", dbname="test")
```

Another alternative is to pass the user/password information in the arguments to the connection function:

```
> con <- TScconnect("ODBC", dbname="test", uid=user, pwd=passwd)
```

This is may be cumbersome to change, and is generally considered to be less secure.

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```
> detach(package:TSodbc)
> detach(package:RODBC)
```

7.4.1 Example ODBC configuration file

Following is an example ODBC configuration file I use in Linux (so the file is in my home directory and called ".odbc.ini") to connect to a remote PostgreSQL server:

[test]

Description	= test DB (Postgresql)
Driver	= Postgresql
Trace	= No
TraceFile	= /tmp/test_odbc.log
Database	= test
Servename	= some.host
UserName	= paul
Password	= mySecret
Port	= 5432
Protocol	= 6.4
ReadOnly	= No
RowVersioning	= No
ShowSystemTables	= No
ShowOidColumn	= No
FakeOidIndex	= No
ConnSettings	=

[ets]

Description	= ets DB (Postgresql)
Driver	= Postgresql
Trace	= No
TraceFile	= /tmp/test_odbc.log
Database	= ets
Servename	= some.host
UserName	= paul
Password	= mySecret
Port	= 5432

```

Protocol          = 6.4
ReadOnly          = No
RowVersioning     = No
ShowSystemTables  = No
ShowOidColumn     = No
FakeOidIndex      = No
ConnSettings      =

```

The above depends on the driver tag “Postgresql” being defined in the file `/etc/odbcinst.ini`, to give the actual driver file location. That file might have something like

```

[PostgreSQL]
Description = PostgreSQL ODBC driver (Unicode version)
Driver      = /usr/lib/x86_64-linux-gnu/odbc/psqlodbcw.so
Setup       = /usr/lib/x86_64-linux-gnu/odbc/libodbcpsqlS.so
Debug      = 0
CommLog     = 1
UsageCount  = 1

```

7.5 TSOacle Connection Details

This package is available on R-forge, but is not being tested, because I do not currently have a server to test it. The code in this section of the vignette is not being run. Please contact the package maintainer (Paul Gilbert) if you have an Oracle server and are willing to test the package.

The Oracle user, password, and hostname should be set in Oracle client configuration file (`tnsnames.ora`) before starting R.

The next small section of code uses `dbConnect` to set up database tables that expected by `TSconnect`.

```

> library("TSOracle")
> con <- dbConnect("Oracle", dbname="test")
> source(system.file("TSsql/CreateTables.TSsql", package = "TSdbi"))
> dbDisconnect(con)

```

Now a TSdbi connection to the database is established.

```

> con <- TSconnect("Oracle", dbname="test")

```

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```

> detach(package:TSOracle)
> detach(package:ROracle)

```


7.6 TSfame Connection Details

I no longer have access to Fame so package *TSfame* is no longer being extensively tested. The code in this section of the vignette is not being run. Please contact the package maintainer (Paul Gilbert) if you have Fame and are willing to test the package.

Beware that the package *fame* may be installed but not functional because the Fame HLI code is not available. A warning will be issued in this case.

Two variants of the Fame connect are available. The first requires a Fame database available on the local system:

```
> con <- TSconnect("fame", dbname="testFame.db")
```

The second requires a Fame server:

```
> con <- TSconnect("fame", dbname="ets /path/to/etsmfacansim.db", "read")
```

where the characters before the space in the dbname string indicate the network name of the server, and the path after the string indicates where the server should find the database.

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```
> detach(package:TSfame)
> detach(package:fame)
```

7.7 TSpadi Connection Details

TSpadi has been largely replaced by other TS* packages and is not longer extensively tested. The code in this section of the vignette is not being run.

TSpadi is a *TSdbi* style interface to a server which extracts data from a backend database. At the Bank of Canada, Fame was the backend database engine. For simplified testing, the backend is a simple set of files, each containing a time series, in a directory which the server can read.

The R client communicates with the server using an RPC based protocol which has been called “padi”, so only the server requires the Fame programs and licences.

TSpadi does not support writing data to a database, so examples below are restricted to an already available database. Alternate approaches for interfacing to Fame are the R packages *fame* and *TSfame*. The later is just a wrapper for the former, to give it a *TSdbi* style interface. Both allow writing to the Fame database, but both also require the Fame HLI programs and licences for the client R machine.

The username and password are not supported, but allowed for compatibility with other packages. The host server should be specified. The dbname may or may not be used, depending on how the server is configured. The user, password, and host can be specified in a single line file (.padi.cfg) in the user’s home directory. The line should be “user password host”. The R code reads this file. Alternatively, the fields can be passed directly to *TSconnect*.

```
> library("TSpadi")
> con <- TSconnect("padi", dbname="ets") # pass user/passwd/host in .padi.cfg
```

Another alternative is

```
> con <- TSconnect("padi", dbname="ets", username=user, password=passwd, host=host)
```

While *TSpadi* does not require Fame on the R clients, there is additional work to set up the padi server and configure it to load databases. Please check additional comments in the last section “TS PADI Server” before deciding about using the *TSpadi* interface (and you may want to contact me, Paul Gilbert). I do actively use this interface, but one of the main purposes of *TSpadi* is to standardize on the *TSdbi* API so that it is easier to move among different database engines.

While it may not be necessary to *detach* packages, the following prevents warnings later about objects being masked:

```
> detach(package:TSpadi)
```

7.8 TS PADI Server

Building a padi server will require some programming effort. The padi interface is getting to be fairly old and, although it still works, some of the underlying code should probably be replaced with a newer approach, perhaps based on something like SOAP.

Code and a description of a prototype of a standard for a Time Series Protocol for Application - Database Interface (TS PADI) may still be available at <http://www.bank-banque-canada.ca/pgilbert>. The code includes a working interface to a Fame database.

8 Appendix B: Underlying Database Structure and Loading Data

More detailed description of the instructions for building the database tables is given in the vignette for the *TSdbi* package. Those instruction show how to build the database using database utilites rather than R, which might be the way a system administrator would build the database.

The database tables are shown in the Table below. The *Meta* table is used for storing meta data about series, such as a description and longer documentation, and also includes an indication of what table the series data is stored in. To retrieve series it is not necessary to know which table the series is in, since this can be found on the *Meta* table. Putting data on the database may require specifying the table, if it cannot be determined from the R representation of the series.

In addition, there will be tables “vintages” and “panels” if those features are used.

Table 1: Data Tables

Table	Contents
Meta	meta data and index to series data tables
A	annual data
Q	quarterly data
M	monthly data
S	semiannual data
W	weekly data
D	daily data
B	business data
U	minutely data
I	irregular data with a date
T	irregular data with a date and time

The following is done with `dbConnect` in place of a `TSconnect`, since they are direct SQL queries and do not require the `TSdbi` structure.

The structure reported reflects the setup that was done previously. These queries are Mysql specific but below is a generic SQL way to do this.

```
> library("TSMysql")
> con <- dbConnect("MySQL", dbname="test")
> dbListTables(con)

[1] "A"      "B"      "D"      "I"      "M"      "Meta" "Q"      "S"      "T"      "U"
[11] "W"

> dbGetQuery(con, "show tables;")

Tables_in_test
1      A
2      B
3      D
4      I
5      M
6      Meta
7      Q
8      S
9      T
10     U
11     W

> dbGetQuery(con, "describe A;")

Field      Type Null Key Default Extra
1  id varchar(40) YES MUL  <NA>
```

```

2 year      int(11) YES MUL    <NA>
3 v         double YES      <NA>

```

```
> dbGetQuery(con, "describe B;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	date	date	YES	MUL	<NA>	
3	period	int(11)	YES	MUL	<NA>	
4	v	double	YES		<NA>	

```
> dbGetQuery(con, "describe D;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	date	date	YES	MUL	<NA>	
3	period	int(11)	YES	MUL	<NA>	
4	v	double	YES		<NA>	

```
> dbGetQuery(con, "describe M;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	year	int(11)	YES	MUL	<NA>	
3	period	int(11)	YES	MUL	<NA>	
4	v	double	YES		<NA>	

```
> dbGetQuery(con, "describe Meta;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	NO	PRI	<NA>	
2	tbl	char(1)	YES	MUL	<NA>	
3	refperiod	varchar(10)	YES		<NA>	
4	description	text	YES		<NA>	
5	documentation	text	YES		<NA>	

```
> dbGetQuery(con, "describe U;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	date	timestamp	NO	MUL	CURRENT_TIMESTAMP	on update CURRENT_TIMESTAMP
3	tz	varchar(4)	YES		<NA>	
4	period	int(11)	YES	MUL	<NA>	
5	v	double	YES		<NA>	

```
> dbGetQuery(con, "describe Q;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	year	int(11)	YES	MUL	<NA>	
3	period	int(11)	YES	MUL	<NA>	
4	v	double	YES		<NA>	

```
> dbGetQuery(con, "describe S;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	year	int(11)	YES	MUL	<NA>	
3	period	int(11)	YES	MUL	<NA>	
4	v	double	YES		<NA>	

```
> dbGetQuery(con, "describe W;")
```

	Field	Type	Null	Key	Default	Extra
1	id	varchar(40)	YES	MUL	<NA>	
2	date	date	YES	MUL	<NA>	
3	period	int(11)	YES	MUL	<NA>	
4	v	double	YES		<NA>	

If schema queries are supported then table information can be obtained in a (almost) generic SQL way. On some systems this will fail because users do not have read privileges on the INFORMATION_SCHEMA table. This does not seem to be an issue in SQLite, but SQLite schema queries are not the same.

```
> dbGetQuery(con, paste(
  "SELECT COLUMN_NAME FROM INFORMATION_SCHEMA.Columns ",
  " WHERE TABLE_SCHEMA='test' AND table_name='A' ;"))
```

	COLUMN_NAME
1	id
2	year
3	v

```
> dbGetQuery(con, paste(
  "SELECT COLUMN_NAME, COLUMN_DEFAULT, COLLATION_NAME, DATA_TYPE,",
  "CHARACTER_SET_NAME, CHARACTER_MAXIMUM_LENGTH, NUMERIC_PRECISION",
  "FROM INFORMATION_SCHEMA.Columns WHERE TABLE_SCHEMA='test' AND table_name='A' ;"))
```

	COLUMN_NAME	COLUMN_DEFAULT	COLLATION_NAME	DATA_TYPE	CHARACTER_SET_NAME
1	id	<NA>	latin1_swedish_ci	varchar	latin1
2	year	<NA>	<NA>	int	<NA>
3	v	<NA>	<NA>	double	<NA>

	CHARACTER_MAXIMUM_LENGTH	NUMERIC_PRECISION
1	40	NA
2	NA	10
3	NA	22

```
> dbGetQuery(con, paste(
  "SELECT COLUMN_NAME, DATA_TYPE, CHARACTER_MAXIMUM_LENGTH, NUMERIC_PRECISION",
  "FROM INFORMATION_SCHEMA.Columns WHERE TABLE_SCHEMA='test' AND table_name='M';"))
```

	COLUMN_NAME	DATA_TYPE	CHARACTER_MAXIMUM_LENGTH	NUMERIC_PRECISION
1	id	varchar	40	NA
2	year	int	NA	10
3	period	int	NA	10
4	v	double	NA	22

9 Appendix C: Examples Using DBI and direct SQL Queries

The following examples are queries using the underlying "DBI" functions. They should not often be needed to access time series, but may be useful to get at more detailed information, or formulate special queries. Typically these queries may be more useful for systems administrators doing database maintenance than they is for end users.

These queries depend on the underlying structure of the database, which should be considered "opaque" from the perspective of a TSdbi user. That is, this structure could be changed without affecting the TSdbi functionality, but the following queries would be affected.

```
> library("TSMysql")
> con <- TSMconnect("MySQL", dbname="test")
> dbGetQuery(con, "SELECT count(*) FROM Meta ;")
```

```
count(*)
1      0
```

```
> dbGetQuery(con, "SELECT max(year) FROM A ;")
```

```
max(year)
1      NA
```

Finally, to disconnect gracefully, one should

```
> dbDisconnect(con)
```

References

- Conway, J., Eddelbuettel, D., Nishiyama, T., Prayaga, S.~K., and Tiffin, N. (2012). *RPostgreSQL: R interface to the PostgreSQL database system*. R package version 0.3-2.
- James, D.~A. (2011). *RSQLite: SQLite interface for R*. R package version 0.11.1.

- James, D.~A. and DebRoy, S. (2012). *RMySQL: R interface to the MySQL database*. R package version 0.9-3.
- R Core Team (2012). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. ISBN 3-900051-07-0.
- R Special Interest Group on Databases (2009). *DBI: R Database Interface*. R package version 0.2-5.
- Ripley, B. and from 1999~to Oct 2002 Michael~Lapsley (2012). *RODBC: ODBC Database Access*. R package version 1.3-6.
- Ryan, J.~A. (2011). *quantmod: Quantitative Financial Modelling Framework*. R package version 0.3-17.
- Ryan, J.~A. and Ulrich, J.~M. (2011). *xts: eXtensible Time Series*. R package version 0.8-2.
- Trapletti, A. and Hornik, K. (2012). *tseries: Time Series Analysis and Computational Finance*. R package version 0.10-28.
- Warnes, G.~R., with contributions~from Ben~Bolker, Gorjanc, G., Grothendieck, G., Korosec, A., Lumley, T., MacQueen, D., Magnusson, A., Rogers, J., and others (2011). *gdata: Various R programming tools for data manipulation*. R package version 2.8.2.
- Zeileis, A. and Grothendieck, G. (2005). zoo: S3 infrastructure for regular and irregular time series. *Journal of Statistical Software*, 14(6):1–27.