

Time Series Database Interface: R PostgreSQL (TSPostgreSQL)

March 6, 2009

1 Introduction

The code from the vignette that generates this guide can be loaded into an editor with `edit(vignette("TSPostgreSQL"))`. This uses the default editor, which can be changed using `options()`. It should be possible to view the pdf version of the guide for this package with `print(vignette("TSPostgreSQL"))`.

WARNING: running these example will overwrite tables in the PostgreSQL "test" database on the server.

The PostgreSQL user, password, and hostname should be set in PostgreSQL information file (.pgpass) before starting R. Alternately, this information can be set with environment variables POSTGRES_USER, POSTGRES_PASSWD and POSTGRES_HOST. (An environment variable POSTGRES_DATABASE can also be set, but "test" is specified below.)

Once R is started, the functions in this package are made available with

```
> library("TSPostgreSQL")
```

This will also load required packages *TSdbi*, *DBI*, *RPostgreSQL*, *methods*, and *tframe*. Some examples below also require *zoo*, and *tseries*.

The next small section of code is necessary to setup database tables that are used in the examples below. It needs to be done only once for a database and might typically be done by an administrator setting up the database, rather than by an end user.

```
> m <- dbDriver("PostgreSQL")
> con <- dbConnect(m, dbname = "test", host = if (!is.null(Sys.getenv("PGHOST"))) Sys.getenv("PGHOST") else "localhost",
> source(system.file("TSsql/CreateTables.TSsql", package = "TSdbi"))
> dbDisconnect(con)
```

More detailed description of the instructions for building the database tables is given in the vignette for the *TSdbi* package. Those instruction show how to build the database using database utilites rather than R, which might be the way a system administrator would build the database.

2 Using the Database - TSdbi Functions

This section gives several simple examples of putting series on and reading them from the database. (If a large number of series are to be loaded into a database, one would typically do this with a batch process using the database program's utilities for loading data.) The first thing to do is to establish a connection to the database:

```
> m <- dbDriver("PostgreSQL")
> con <- TSconnect(m, dbname = "test")
```

TSconnect uses *dbConnect* from the *DBI* package, but checks that the database has expected tables, and checks for additional features. (It cannot be used before the tables are created, as done in the previous section.)

This puts a series called *vec* on the database and then reads it back

```
> z <- ts(rnorm(10), start = c(1990, 1), frequency = 1)
> seriesNames(z) <- "vec"
> if (TSexists("vec", con)) TSdelete("vec", con)
> TSput(z, con)
> z <- TSget("vec", con)
```

If the series is printed it is seen to be a "ts" time series with some extra attributes.

TSput fails if the series already exists on the *con*, so the above example checks and deletes the series if it already exists. *TSreplace* does not fail if the series does not yet exist, so examples below use it instead. Several plots below show original data and the data retrieved after it is written to the database. One is added to the original data so that both lines are visible.

And now more examples:

```
> z <- ts(matrix(rnorm(20), 10, 2), start = c(1990, 1), frequency = 1)
> seriesNames(z) <- c("matc1", "matc2")
> TSreplace(z, con)
```

```
[1] TRUE
```

```
> TSget("matc1", con)
```

```
Time Series:
```

```
Start = 1990
```

```
End = 1999
```

```
Frequency = 1
```

1	2	3	4	5	6
-0.39547850	0.84223662	-1.41937916	0.28428632	-1.62880715	0.98090499
7	8	9	10		
0.31455997	-0.83015434	0.07511815	-0.35556201		

```
attr("seriesNames")
```

```
[1] matc1
```

```

attr("TSrefperiod")
[1] NA
attr("TSMeta")
An object of class "IJTSMeta"
Slot "TSdescription":
[1] NA

Slot "TSdoc":
[1] NA

Slot "TSlabel":
[1] NA

Slot "serIDs":
[1] "matc1"

Slot "conType":
[1] "TSPostgreSQLConnection"
attr("package")
[1] "TSPostgreSQL"

Slot "DateStamp":
[1] NA

Slot "dbname":
[1] "test"

Slot "hasVintages":
[1] FALSE

Slot "hasPanels":
[1] FALSE

> TSget("matc2", con)

Time Series:
Start = 1990
End = 1999
Frequency = 1
      1          2          3          4          5          6
-0.60173530 -1.20505336 -0.68241656 -0.08106582 -0.86782853  0.16346720
      7          8          9         10
 0.01103349  0.53813930 -0.37401290 -0.34712384
attr("seriesNames")
[1] matc2
attr("TSrefperiod")

```

```

[1] NA
attr(,"TSMeta")
An object of class "IJTSMeta"
Slot "TSdescription":
[1] NA

Slot "TSdoc":
[1] NA

Slot "TSlabel":
[1] NA

Slot "serIDs":
[1] "matc2"

Slot "conType":
[1] "TSPostgreSQLConnection"
attr(,"package")
[1] "TSPostgreSQL"

Slot "DateStamp":
[1] NA

Slot "dbname":
[1] "test"

Slot "hasVintages":
[1] FALSE

Slot "hasPanels":
[1] FALSE

> TSget(c("matc1", "matc2"), con)

Time Series:
Start = 1990
End = 1999
Frequency = 1
      matc1      matc2
1990 -0.39547850 -0.60173530
1991  0.84223662 -1.20505336
1992 -1.41937916 -0.68241656
1993  0.28428632 -0.08106582
1994 -1.62880715 -0.86782853
1995  0.98090499  0.16346720
1996  0.31455997  0.01103349

```

```

1997 -0.83015434  0.53813930
1998  0.07511815 -0.37401290
1999 -0.35556201 -0.34712384
attr(,"TSrefperiod")
[1] NA NA
attr(,"TSMeta")
An object of class "TSmeta"
Slot "TSdescription":
[1] NA

Slot "TSdoc":
[1] NA

Slot "TSlabel":
[1] NA

Slot "serIDs":
[1] "matc1" "matc2"

Slot "conType":
[1] "TSPostgreSQLConnection"
attr(,"package")
[1] "TSPostgreSQL"

Slot "DateStamp":
[1] NA



Slot "dbname":
[1] "test"

Slot "hasVintages":
[1] FALSE

Slot "hasPanels":
[1] FALSE

> tfplot(z + 1, TSget(c("matc1", "matc2"), con), lty = c("solid",
    "dashed"), col = c("black", "red"))

```



```
> z <- ts(matrix(rnorm(20), 10, 2), start = c(1990, 1), frequency = 4)
> seriesNames(z) <- c("matc1", "matc2")
> TSreplace(z, con)
```

```
[1] TRUE
```

```
> TSget(c("matc1", "matc2"), con)
```

	matc1	matc2
1990 Q1	0.02864730	1.00866982
1990 Q2	-1.17601388	-0.01187087
1990 Q3	-0.15805214	1.37428591
1990 Q4	-0.08954945	-1.20612684
1991 Q1	-0.63897393	-2.18627586
1991 Q2	0.38549596	1.09726645
1991 Q3	-1.35838282	0.91629434
1991 Q4	1.17899015	0.06327174
1992 Q1	-0.48293866	2.13937633
1992 Q2	1.04597473	1.03478856

```
attr("TSrefperiod")
[1] NA NA
attr("TSmeta")
An object of class "TSmeta"
```

```

Slot "TSdescription":
[1] NA

Slot "TSdoc":
[1] NA

Slot "TSlabel":
[1] NA

Slot "serIDs":
[1] "matc1" "matc2"

Slot "conType":
[1] "TSPostgreSQLConnection"
attr(,"package")
[1] "TSPostgreSQL"

Slot "DateStamp":
[1] NA



Slot "dbname":
[1] "test"

Slot "hasVintages":
[1] FALSE

Slot "hasPanels":
[1] FALSE

> tfplot(z + 1, TSget(c("matc1", "matc2"), con), lty = c("solid",
    "dashed"), col = c("black", "red"))



```



```
> z <- ts(matrix(rnorm(200), 100, 2), start = c(1995, 1), frequency = 12)
> seriesNames(z) <- c("mat2c1", "mat2c2")
> TSreplace(z, con)

[1] TRUE

> tfplot(z + 1, TSget(c("mat2c1", "mat2c2"), con), lty = c("solid",
  "dashed"), col = c("black", "red"))
```

The following extract information about the series from the database, although not much information has been added for these examples.

```
> TSmeta("mat2c1", con)
> TSmeta("vec", con)
> TSdates("vec", con)
> TSdescription("vec", con)
> TSdoc("vec", con)
```

Below are examples that make more use of *TSdescription* and *codeTSdoc*. Often it is convenient to set the default connection:

```
> options(TSconnection = con)
```

and then the *con* specification can be omitted from the function calls unless another connection is needed. The *con* can still be specified, and some examples below do specify it, just to illustrate the alternative syntax.

```
> z <- TSget("mat2c1")
> TSmeta("mat2c1")
```

```
An object of class "TSmeta"
Slot "TSdescription":
[1] "NA"
```

```

Slot "TSdoc":
[1] "NA"

Slot "TSlabel":
[1] NA

Slot "serIDs":
[1] "mat2c1"

Slot "conType":
[1] "TSPostgreSQLConnection"
attr(,"package")
[1] "TSPostgreSQL"

Slot "DateStamp":
[1] NA

Slot "dbname":
[1] "test"

Slot "hasVintages":
[1] FALSE

Slot "hasPanels":
[1] FALSE

```

Data documentation can be in two forms, a description specified by *TSdescription* or longer documentation specified by *TSdoc*. These can be added to the time series object, in which case they will be written to the database when *TSput* or *TSreplace* is used to put the series on the database. Alternatively, they can be specified as arguments to *TSput* or *TSreplace*. The description or documentation will be retrieved as part of the series object with *TSget* only if this is specified with the logical arguments *TSdescription* and *TSdoc*. They can also be retrieved directly from the database with the functions *TSdescription* and *TSdoc*.

```

> z <- ts(matrix(rnorm(10), 10, 1), start = c(1990, 1), frequency = 1)
> TSreplace(z, serIDs = "Series1", con)

[1] TRUE

> zz <- TSget("Series1", con)
> TSreplace(z, serIDs = "Series1", con, TSdescription = "short rnorm series",
  TSdoc = "Series created as an example in the vignette.")

[1] TRUE

```

```

> zz <- TSget("Series1", con, TSdescription = TRUE, TSdoc = TRUE)
> start(zz)

[1] 1990    1

> end(zz)

[1] 1999    1

> TSdescription(zz)

[1] "short rnorm series"

> TSdoc(zz)

[1] "Series created as an example in the vignette."

> TSdescription("Series1", con)

[1] "short rnorm series"

> TSdoc("Series1", con)

[1] "Series created as an example in the vignette."

> z <- ts(rnorm(10), start = c(1990, 1), frequency = 1)
> seriesNames(z) <- "vec"
> TSreplace(z, con)



[1] TRUE

> zz <- TSget("vec", con)
> z <- ts(matrix(rnorm(20), 10, 2), start = c(1990, 1), frequency = 1)
> seriesNames(z) <- c("matc1", "matc2")
> TSreplace(z, con)

[1] TRUE

> tfplot(z + 1, TSget(c("matc1", "matc2"), con), lty = c("solid",
  "dashed"), col = c("black", "red"))



```



```
> z <- ts(matrix(rnorm(20), 10, 2), start = c(1990, 1), frequency = 4)
> seriesNames(z) <- c("matc1", "matc2")
> TSreplace(z, con)

[1] TRUE



> tfplot(z + 1, TSget(c("matc1", "matc2"), con), lty = c("solid",
  "dashed"), col = c("black", "red"))
```



```
> z <- ts(matrix(rnorm(200), 100, 2), start = c(1995, 1), frequency = 12)
> seriesNames(z) <- c("mat2c1", "mat2c2")
> TSreplace(z, con)

[1] TRUE

> tfplot(z + 1, TSget(c("mat2c1", "mat2c2"), con), lty = c("solid",
  "dashed"), col = c("black", "red"))
```





The following examples use dates and times which are not handled by *ts*, so the *zoo* time representation is used.

```
> require("zoo")
> z <- zoo(matrix(rnorm(200), 100, 2), as.Date("1990-01-01") +
  0:99)
> seriesNames(z) <- c("zooc1", "zooc2")
> TSreplace(z, con, Table = "D")

[1] TRUE



> tfplot(z + 1, TSget(c("zooc1", "zooc2"), con), lty = c("solid",
  "dashed"), col = c("black", "red"))
```



```
> z <- zoo(matrix(rnorm(200), 100, 2), as.Date("1990-01-01") +
  0:99 * 7)
> seriesNames(z) <- c("zooWc1", "zooWc2")
> TSreplace(z, con, Table = "W")

[1] TRUE

> tfplot(z + 1, TSget(c("zooWc1", "zooWc2"), con), col = c("black",
  "red"), lty = c("dashed", "solid"))
```



```
> dbDisconnect(con)
```

3 Examples Using Web Data

This section illustrates fetching data from a web server and loading it into the database. This would be a very slow way to load a database, but provides examples of different kinds of time series data. The fetching is done with *TShistQuote* which provides a wrapper for *get.hist.quote* from package *tseries* to give syntax consistent with the *TSdbi*.

Fetching data may fail due to lack of an Internet connection or delays.

First establish a connection to the database where data will be saved:

```
> con <- TSconnect("PostgreSQL", dbname = "test")
```

Now connect to the web server and fetch data:

```
> require("TShistQuote")
> Yahoo <- TSconnect("histQuote", dbname = "yahoo")
> x <- TSget("^gspc", quote = "Close", con = Yahoo)
> plot(x)
> tfplot(x)
> TSrefperiod(x)
```



```

[1] "Close"
> TSdescription(x)
[1] "^gspc Close from yahoo"
> TSdoc(x)
[1] "^gspc Close from yahoo retrieved 2009-03-06 22:14:58"
> TSlabel(x)
[1] "^gspc Close"

```

Then write the data to the local server, specifying table B for business day data (using `TSreplace` in case the series is already there from running this example previously):

```

> TSreplace(x, serIDs = "gspc", Table = "B", con = con)
[1] TRUE


```

and check the saved version:

```

> TSrefperiod(TSget(serIDs = "gspc", con = con))
[1] "Close"
> TSdescription("gspc", con = con)
[1] "^gspc Close from yahoo"
> TSdoc("gspc", con = con)
[1] "^gspc Close from yahoo retrieved 2009-03-06 22:14:58"
> TSlabel("gspc", con = con)
[1] NA
> tfplot(TSget(serIDs = "gspc", con = con))

```



```
> x <- TSget("ibm", quote = c("Close", "Vol"), con = Yahoo)
> TSreplace(x, serIDs = c("ibm.Cl", "ibm.Vol"), con = con, Table = "B",
  TSdescription. = c("IBM Close", "IBM Volume"), TSdoc. = paste(c("IBM Close retrieved on ", Sys.Date()),
    "IBM Volume retrieved on "), Sys.Date()))

[1] TRUE



> z <- TSget(serIDs = c("ibm.Cl", "ibm.Vol"), TSdescription = TRUE,
  TSdoc = TRUE, con = con)
> TSdescription(z)

[1] "IBM Close" "IBM Volume"

> TSdoc(z)

[1] "IBM Close retrieved on 2009-03-06"
[2] "IBM Volume retrieved on 2009-03-06"

> tfplot(z, xlab = TSdoc(z), Title = TSdescription(z))
> tfplot(z, Title = "IBM", start = "2007-01-01")
```




Oanda has maximum of 500 days, so the start date is specified here so as to not exceed that.

```
> Oanda <- TSconnect("histQuote", dbname = "oanda")
> x <- TSget("EUR/USD", start = Sys.Date() - 495, con = Oanda)
> TSreplace(x, serIDs = "EUR/USD", Table = "D", con = con)
```

```
[1] TRUE
```

Then check the saved version:

```
> z <- TSget(serIDs = "EUR/USD", TSlabel = TRUE, TSdescription = TRUE,
  con = con)
> tfplot(z, Title = TSdescription(z), ylab = TSlabel(z))
> tfplot(z, Title = "EUR/USD", start = "2007-01-01")
> tfplot(z, Title = "EUR/USD", start = "2007-03-01")
> tfplot(z, Title = "EUR/USD", start = Sys.Date() - 14, end = Sys.Date(),
  xlab = format(Sys.Date(), "%Y"))
```



```
> dbDisconnect(con)
> dbDisconnect(Yahoo)
> dbDisconnect(Oanda)
```

3.1 Examples Using TSdbi with ets

These examples use a database called "ets" which is available at the Bank of Canada. This set of examples illustrates how the programs might be used if a larger database is available. Typically a large database would be installed using database scripts directly rather than from R with *TSput* or *TSreplace*.

The following are wrapped in *if (!inherits(conets, "try-error"))* so that the vignette will build even when the database is not available. This seems to require an explicit call to *print()*, but that is not usually needed to display results below. Another artifact of this is that results printed in the if block do not display until the end of the block.

```
> m <- dbDriver("PostgreSQL")
> conets <- try(TSconnect(m, dbname = "ets"))
> if (!inherits(conets, "try-error")) {
  options(TSconnection = conets)
  print(TSmeta("M.SDR.CCUSMA02.ST"))
  EXCH.IDs <- t(matrix(c("M.SDR.CCUSMA02.ST", "SDR/USD exchange rate",
```

```

      "M.CAN.CCUSMA02.ST", "CAN/USD exchange rate", "M.MEX.CCUSMA02.ST",
      "MEX/USD exchange rate", "M.JPN.CCUSMA02.ST", "JPN/USD exchange rate",
      "M.EMU.CCUSMA02.ST", "Euro/USD exchange rate", "M.OTO.CCUSMA02.ST",
      "OECD /USD exchange rate", "M.G7M.CCUSMA02.ST", "G7 /USD exchange rate",
      "M.E15.CCUSMA02.ST", "Euro 15. /USD exchange rate"),
      2, 8))
print(TSdates(EXCH.IDs[, 1]))
z <- TSdates(EXCH.IDs[, 1])
print(start(z))
print(end(z))
tfplot(TSget(serIDs = "V122646", conets))
}

An object of class "IJTSmeta"
Slot "TSdescription":
[1] "Special Drawing Right---Currency Conversions/US$ exchange rate/Average of daily rates/M

Slot "TSdoc":
[1] "Special Drawing Right---Currency Conversions/US$ exchange rate/Average of daily rates/M

Slot "TSlabel":
[1] NA

Slot "serIDs":
[1] "M.SDR.CCUSMA02.ST"

Slot "conType":
[1] "TSPostgreSQLConnection"
attr(,"package")
[1] "TSPostgreSQL"

Slot "DateStamp":
[1] NA

Slot "dbname":
[1] "ets"

Slot "hasVintages":
[1] FALSE

Slot "hasPanels":
[1] FALSE

[,1]
[1,] "M.SDR.CCUSMA02.ST from 1960 1 to 2007 9 M NA "
[2,] "M.CAN.CCUSMA02.ST from 1960 1 to 2007 9 M NA "

```

```

[3,] "M.MEX.CCUSMA02.ST from 1963 1 to 2007 9 M    NA    "
[4,] "M.JPN.CCUSMA02.ST from 1960 1 to 2007 9 M    NA    "
[5,] "M.EMU.CCUSMA02.ST from 1979 1 to 2007 9 M    NA    "
[6,] "M.OTO.CCUSMA02.ST    not available"
[7,] "M.G7M.CCUSMA02.ST    not available"
[8,] "M.E15.CCUSMA02.ST    not available"
[[1]]
[1] 1960      1

[[2]]
[1] 1960      1

[[3]]
[1] 1963      1

[[4]]
[1] 1960      1

[[5]]
[1] 1979      1

[[6]]
[1] NA

[[7]]
[1] NA

[[8]]
[1] NA

[[1]]
[1] 2007      9

[[2]]
[1] 2007      9

[[3]]
[1] 2007      9

[[4]]
[1] 2007      9

[[5]]
[1] 2007      9

[[6]]

```


```
[1] NA
```

```
[[7]]
```


```
[1] NA
```

```
[[8]]
```


```
[1] NA
```




```
> if (!inherits(conets, "try-error")) {  
  print(TSdescription(TSget("V122646", TSdescription = TRUE)))  
  print(TSdescription("V122646"))  
  print(TSdoc(TSget("V122646", TSdoc = TRUE)))  
  print(TSdoc("V122646"))  
  tfplot(TSget("V122646", names = "V122646", conets))  
}  
  
[1] "Total short-term business credit, Seasonally adjusted, average of month-end"  
[1] "Total short-term business credit, Seasonally adjusted, average of month-end"  
[1] "Same as B171"  
[1] "Same as B171"
```





```
> if (!inherits(conets, "try-error")) {  
  z <- TSget("V122646", TSdescription = TRUE)  
  tfplot(z, Title = strsplit(TSdescription(z), ","))  
}
```

```
> if (!inherits(conets, "try-error")) {  
  z <- TSget("SDSP500", TSdescription = TRUE)  
  tfplot(z, Title = TSdescription(z))  
  plot(z)  
}
```




```
> if (!inherits(conets, "try-error")) {  
  z <- TSget(c("DSP500", "SDSP500"), TSdescription = TRUE)  
  tfplot(z, xlab = TSdescription(z))  
}
```




```
> if (!inherits(conets, "try-error")) {
  plot(z)
}
```

z



```
> if (!inherits(conets, "try-error")) {  
  ETSgdp <- annualizedGrowth(aggregate(TSget("V1992067"), nfrequency = 4,  
    FUN = mean), lag = 4, names = "GDP Y/Y Growth")  
  tfplot(ETSgdp)  
}
```



```
> if (!inherits(conets, "try-error")) {
  dbDisconnect(options()$TSconnection)
  options(TSconnection = NULL)
}
```

4 Examples Using DBI and direct SQL Queries

The following examples are queries using the underlying "DBI" functions. They should not often be needed to access time series, but may be useful to get at more detailed information, or formulate special queries.

```
> m <- dbDriver("PostgreSQL")
> con <- TSconnect(m, dbname = "test")
> options(TSconnection = con)

> dbListTables(con)

[1] "usarrests" "meta"      "a"          "b"          "d"          "m"
[7] "u"         "q"         "s"          "w"          "i"          "t"
```

If schema queries are supported then table information can be found in a generic SQL way, but on some systems this will fail because users do not have

read privileges on the INFORMATION_SCHEMA table, so the following are wrapped in *try()*. (SQLite does not seem to support this at all.)

```
> try(dbGetQuery(con, paste("SELECT COLUMN_NAME FROM INFORMATION_SCHEMA.Columns ",
    " WHERE TABLE_SCHEMA='test' AND table_name='A' ;")))
```

data frame with 0 columns and 0 rows

```
> try(dbGetQuery(con, paste("SELECT COLUMN_NAME, COLUMN_DEFAULT, COLLATION_NAME, DATA_TYPE,
    "CHARACTER_SET_NAME, CHARACTER_MAXIMUM_LENGTH, NUMERIC_PRECISION",
    "FROM INFORMATION_SCHEMA.Columns WHERE TABLE_SCHEMA='test' AND table_name='A' ;")))
```

data frame with 0 columns and 0 rows

```
> try(dbGetQuery(con, paste("SELECT COLUMN_NAME, DATA_TYPE, CHARACTER_MAXIMUM_LENGTH, NUMERIC_PRECISION,
    "FROM INFORMATION_SCHEMA.Columns WHERE TABLE_SCHEMA='test' AND table_name='M' ;")))
```

data frame with 0 columns and 0 rows

Finally, to disconnect gracefully, one should

```
> dbDisconnect(con)
> dbDisconnect(options())$TSconnection
> options(TSconnection = NULL)
```